

Tim Hua

Independent AI Safety Researcher. Formerly on John List's economics team at Walmart
timhua0@gmail.com Website: timhua.me

EDUCATION & SKILLS

Middlebury College

BA, *summa cum laude*, Economics (GPA 3.97/4.00)

Middlebury, VT

May 2023

D.K. Smith Economics Prize for best thesis (*Fox News's Effect on Social and Moral Preferences*).

- Thesis cited by Harvard economist Benjamin Enke in [Moral Boundaries](#)

Harvard College

Visiting Undergraduate Student (GPA: 4.0/4.0)

Cambridge, MA

2021 - 2022 School Year

- Took courses economics theory, probability and statistics, real analysis, etc.

General skills: Pytorch/Python, R, SQL, Stata.

AI SAFETY RESEARCH

Neel Nanda MATS Training Program

Scholar

Remote

Apr 2025 - Current

- Accepted (with funding) to Neel's training program in mechanistic interpretability.

Tyler Tracy's MARS Mentee in AI Control

Researcher

Berkeley, CA

Jan 2025 - Current

- Working with Tyler Tracy (Redwood) on evaluating and formalizing new control protocols.
- Modeling how to optimally allocate budget between human auditing, an expensive and accurate monitor, and a cheaper but less accurate monitor.
- Presented poster and gave a lightning talk at Control Conf.

Alice Rigg's AI Safety Camp Mentee in Mechanistic Interpretability

Researcher

Remote

Jan 2025 - Current

- Working with Alice Rigg on mechanistic interpretability studying the functional role of SAE error terms.

Miscellaneous projects

- BlueDot Economics of Transformative AI Course Facilitator (Feb 2025 Cohort)
- Wrote a critique of the bias in bios experiment in Sparse Feature Circuits on [alignment forum](#).
- Conceived and explored a control protocol where the untrusted model writes a plan and delegates the execution to a trusted model. Writeup and [pilot results here](#).
- Interpreting obfuscating jailbreaks together with Veniamin Veselovsky, some [results are here](#).

PROFESSIONAL EXPERIENCES

Walmart Economics Team

Manager, Economist

New York, NY

Intern in summer 2022, full time July 2023 – Aug 2024

Worked on monetary incentives for drivers on Walmart's last mile delivery platform.

- Designed and executed numerous experiments, some with multiple treatment arms.
- Built ML models using causal random forests and other double ML methods, ran hyperparameter sweeps with out-of-sample off-policy validation following [Treatment Allocation Under Uncertain Costs](#)
- Wrote R code that conducts synthetic difference in differences estimation and calculates conformal inference-based confidence in parallel. Code reused by others on my team.
- Presented work to Walmart chief economist John List, [fifth ranked economist in the world](#)
- Translated technical findings and frameworks for non-technical partner business teams

As an intern, created pricing zones that minimized customers' cross-zone shopping by constructing a weighted graph of Walmart stores and partitioning the graph.

- Pricing zones are still used at Walmart today

Harvard Economics Department

Remote

Course Assistant for Professor Edward Glaeser

June 2022 – Dec 2022

- Wrote sections in the textbook of Econ 1011a: Intermedia Microeconomics: Advanced.

Harvard Kennedy School

Cambridge, MA

Research Assistant for Professor Desmond Ang

March 2022 – July 2022

- Scraped tvtropes.org for an exploratory project investigating stereotypes in films.

Brookings Institution

Remote

Research Assistant for Professor Carol Graham

June 2021 – May 2022

- Coauthor on [paper](#) examining mental and physical ill-being during the COVID-19 pandemic.

Middlebury Economics Department

Middlebury, VT

Research Assistant for Professor Marin Abel

September 2020 - March 2021

- Investigate how different compensation schemes could help close the behavior-intention gap.